

COSC 6397

Big Data Analytics

Hadoop 2 and YARN

Edgar Gabriel
Spring 2015

UNIVERSITY of **HOUSTON**

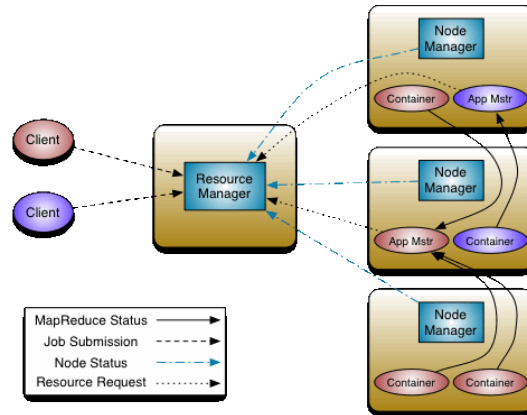
Hadoop 2

- Major differences between MapReduce (v1) before and after hadoop-0.23 => MapReduce 2.0 (MRv2) or YARN.
- Split up the two major functionalities of JobTracker:
 - resource management and
 - job scheduling/monitoring
- Use a global ResourceManager (*RM*) and per-application ApplicationMaster (*AM*)
- An application can be DAG of jobs.
- ResourceManager is the ultimate authority to arbitrate resources among all the applications in the system
- Per-application ApplicationMaster is a framework specific library to negotiate resources from the ResourceManager and working with the NodeManager(s) to execute and monitor the tasks.

<http://hadoop.apache.org/docs/current/hadoop-yarn/hadoop-yarn-site/YARN.html>

UNIVERSITY of **HOUSTON**

YARN architecture



<http://hadoop.apache.org/docs/current/hadoop-yarn/hadoop-yarn-site/YARN.html>

UNIVERSITY of HOUSTON

YARN Platform Benefits

- Deployment
 - YARN provides a seamless vehicle to deploy your software to an enterprise Hadoop cluster
- Fault Tolerance
 - YARN 'handles' (detects, notifies, and provides default actions) for HW, OS, JVM failure tolerance
 - YARN provides plugins for the app to define failure behavior
- Scheduling (incorporating Data Locality)
 - YARN utilizes HDFS to schedule app processing where the data lives
 - YARN ensures that your apps finish in the SLA expected by your customers

Slide based on a lecture by Arun Murthy and Bob Page:
"Developing YARN Native applications"

UNIVERSITY of HOUSTON

Application Categories

Type	Definition	Examples
Framework / Engine	Provides platform capabilities to enable data services and applications	Twil, Reef, Tez, MapReduce, Spark
Service	An application that runs continuously	Storm, HBase, Memcached, etc
YARN App	A <i>temporal job</i> submitted to YARN	MapReduce job, spark job etc.

Slide based on a lecture by Arun Murthy and Bob Page:
"Developing YARN Native applications"

UNIVERSITY of HOUSTON

Yarn Concepts: Container

- Basic unit of allocation
 - Based on capability: e.g, memory, CPU etc
 - Replaces the fixed map/reduce slots from Hadoop 1
- Allows for fine grained resources allocation
 - Capability, Host, Rack, Priority etc.
- `ContainerLaunchContext`: object describing resources necessary to launch the Container
 - Local Resources: Resources needed to execute container application
 - Environment variables: e.g. CLASSPATH
 - Command to execute

Slide based on a lecture by Arun Murthy and Bob Page:
"Developing YARN Native applications"

UNIVERSITY of HOUSTON

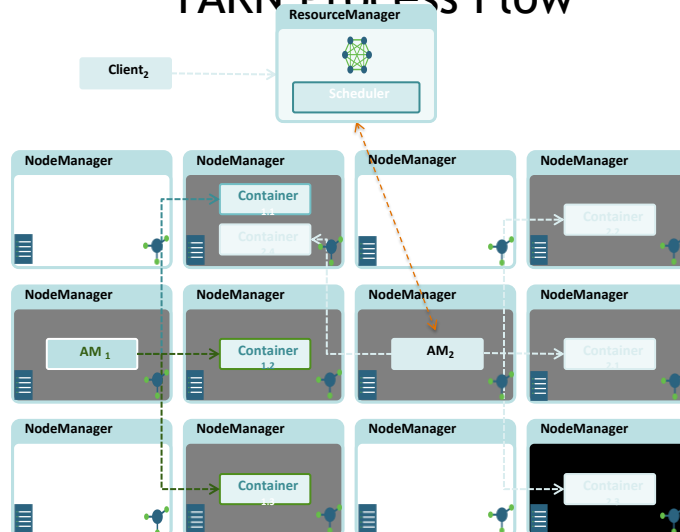
YARN Terminology

- Resource Manager
 - Central agent
 - Allocates & manages cluster resources
- Node Manager
 - Manages, monitors and enforces node resource allocation
 - Manages life cycle of containers
- User application
 - Client: submits application
 - Application Master (AM): manages application lifecycle and task scheduling
 - Container: executes application logic

Slide based on a lecture by Arun Murthy and Bob Page:
"Developing YARN Native applications"

UNIVERSITY of HOUSTON

YARN Process Flow

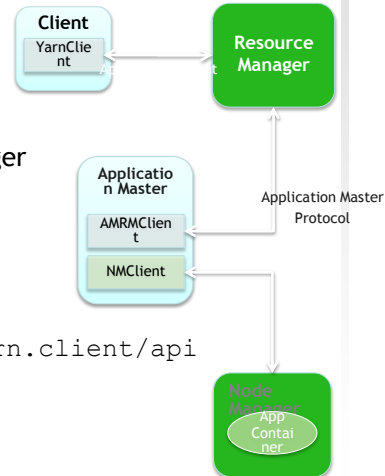


Slide based on a lecture by Arun Murthy and Bob Page:
"Developing YARN Native applications"

UNIVERSITY of HOUSTON

YARN APIs

- Three protocols
 - Client to ResourceManager
 - Application submission
 - ApplicationMaster to ResourceManager
 - Container allocation
 - ApplicationMaster to NodeManager
 - Container launch
- Use client libraries
 - Package `org.apache.hadoop.yarn.client/api`



UNIVERSITY of HOUSTON

YARN Implementation Outline

1. Write a client
 1. Submit the application
 2. Monitor application
2. Write an ApplicationMaster
3. Get containers and run application

Slide based on a lecture by Arun Murthy and Bob Page:
"Developing YARN Native applications"

UNIVERSITY of HOUSTON

Write a client

- Client initializes and starts a YarnClient.

```
YarnClient yarnClient = YarnClient.createYarnClient();
```

- Once a client is set up, the client needs to create an application, and get its application id.

```
YarnClientApplication app = yarnClient.createApplication();
GetNewApplicationResponse appResponse =
app.getNewApplicationResponse();
```

- Response from YarnClientApplication contains information about the cluster (e.g. minimum/maximum resource capabilities)
 - required for correctly setting specifications of the container for ApplicationMaster

UNIVERSITY of **HOUSTON**

Write a client (II)

- Setup the ApplicationSubmissionContext which defines all the information needed by the RM to launch the AM. A client needs to set the following into the context:
 - Application info: id, name
 - Queue, priority info: Queue to which the application will be submitted, the priority to be assigned for the application.
 - User: The user submitting the application
- ContainerLaunchContext: The information defining the container in which the AM will be launched and run.

UNIVERSITY of **HOUSTON**

```
// set the application submission context
appContext = app.getApplicationSubmissionContext();
appContext.setKeepContainersAcrossApplicationAttempts(keepContainers);
appContext.setApplicationName(appName);

// Set up resource type requirements
Resource capability = Resource.newInstance(amMemory, amVCores);
appContext.setResource(capability);

// Set up the container launch context for the application master
ContainerLaunchContext amContainer =
ContainerLaunchContext.newInstance( localResources, env, commands,
null, null, null);
appContext.setAMContainerSpec(amContainer);

Priority pri = Priority.newInstance(amPriority);
appContext.setPriority(pri);
appContext.setQueue(amQueue);

// Submit the application to the applications manager
yarnClient.submitApplication(appContext);
```

UNIVERSITY of **HOUSTON**

Client monitoring

- Client communicates with RM and requests a report of the application via the `getApplicationReport()` method of `YarnClient`:
 - *General application information*: Application id, queue to which the application was submitted, user who submitted the application and the start time for the application.
 - *ApplicationMaster details*: the host on which the AM is running, the rpc port (if any) on which it is listening for requests from clients and a token that the client needs to communicate with the AM.
 - *Application tracking information*: tracking url
 - *Application status*: The state of the application as seen by the `ResourceManager`

UNIVERSITY of **HOUSTON**

Writing ApplicationMaster

- The AM is the actual owner of the job
- Launched by the ResourceManager
- Client provides information about resources required to oversee and complete the job
- AM is launched within a container that may (likely will) be sharing a physical host with other containers
- Several parameters are made available to AM via the environment, e.g.
 - ContainerId for the AM container,
 - application submission time
 - details about the NodeManager host running the AM
- All interactions with the RM require an ApplicationAttemptId

UNIVERSITY of **HOUSTON**

Writing ApplicationMaster

- AM starts two clients:
 - Communication with ResourceManager,
 - Handling NodeManager
- AM emits regular heartbeat signal to RM
- AM calculates number of containers needed based task requirements and requests the container from RM
- AM has to establish the following information for job submission(using `setupContainerAskForRM()`):
 - *Resource capability*: value is defined in MB and has to less than the max capability of the cluster and an exact multiple of the min capability.
 - *Priority*: different priorities can be requested for different sets (e.g. for mapper and reducer)

UNIVERSITY of **HOUSTON**

Writing ApplicationMaster

- After container allocation requests have been sent by the application manager, containers will be launched asynchronously

```
public void onContainersAllocated(List<Container>
                                allocatedContainers) {
    numAllocatedContainers.addAndGet(allocatedContainers.size());
    for (Container allocatedContainer : allocatedContainers) {
        LaunchContainerRunnable runnableLaunchContainer = new
            LaunchContainerRunnable(allocatedContainer,
                                    containerListener);
        Thread launchThread = new Thread(runnableLaunchContainer);
        // launch and start the container on a separate thread to
        // keep the main thread unblocked as all containers may not
        // be allocated at one go.
        launchThreads.add(launchThread); launchThread.start(); } }
```

UNIVERSITY of HOUSTON

Unmanaged mode for ApplicationMaster

- Run the Applicationmaster on developmental machine rather than cluster
 - No submission needed
 - Easier to debug
 - Use `hadoop-yarn-application-unmanaged-am-launcher`

```
$ bin/hadoop jar hadoop-yarn-applications-unmanaged-am-launcher.jar
Client \
-jar my-application-master.jar \
-cmd 'java MyApplicationMaster <args>'
```

UNIVERSITY of HOUSTON

Non MapReduce Applications

- Distributed Shell
- Impala
- Apache Giraph
- Spark
- ...

UNIVERSITY of **HOUSTON**

MapReduce on Yarn

- Hadoop includes a MapReduce Application master to manage MR jobs
- Each Mapreduce job is an instance of a YARN application
- Shuffle is an auxiliary service that runs in the NodeManager JVM

UNIVERSITY of **HOUSTON**

Fault Tolerance of MR jobs

- Task/Container failure
 - MRAppMaster will reexecute tasks that execute with an exception or stop responding (4 times by default)
 - Applications with too many failed tasks considered failed
- Application Master failure
 - If AM stops sending heartbeats, ResourceManager will attempt to start a new master (2 times by default)
- NodeManager
 - If NM stops sending heartbeats to RM, node will be removed from list of available nodes
 - Tasks on that node be treated as failed ab AM
- Resource Manager: single point of failure if not configured in High Availability Mode